Machine Learning Engineer

Company: Saransh AI Labs
Location: Hyderabad, India (hybrid)
Employment type: Full-time

About us

Saransh AI builds language and vision products for banks and insurers. Our models read documents, route customer messages and check claim photos for millions of customers.

About the role

We are hiring a Machine Learning Engineer to design, train and deploy models that run in production. You will own models from the first experiment to monitoring after release.

Responsibilities

- Train deep learning models for natural language processing and computer vision
- Turn research code into reliable services
- Deploy models with Docker and Kubernetes on AWS
- Maintain CI/CD pipelines for training and deployment
- Track experiments and monitor models in production

Requirements

- 3+ years of experience in machine learning engineering
- Strong Python with PyTorch or TensorFlow
- scikit-learn, NumPy and Pandas
- Docker and Kubernetes
- A solid background in statistics
- Comfortable working on Linux
- Master's degree or bachelor's degree in computer science or a related field

Nice to have

- GCP or Azure
- Spark for large datasets
- Published research or open-source work

What we offer

- Health insurance and a yearly GPU budget for side projects
- Conference travel once a year
- Flexible working hours

How to apply

Send your resume and a short description of a model you took to production.